



COEN 241

Introduction to Cloud Computing

Lecture 5 - OS Virtualization





Lecture 4 Recap

- Advanced Virtualization Concepts II
 - Memory Virtualization
 - Paravirtualization
- Vagrant

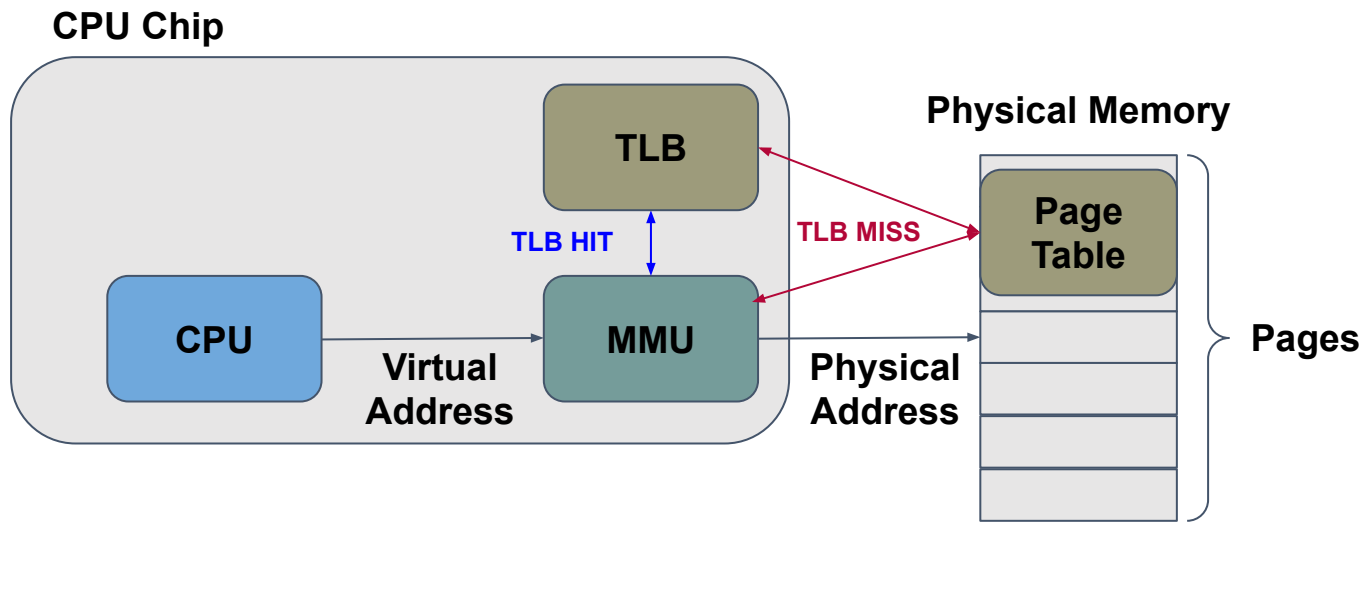


Another Virtualization Challenge: RAM

- RAM access patterns are surprisingly complex
- History: CPU's address pins indicate word to read/write
 - e.g., MOS 6502 has 16 address wires, thus 64kB RAM (2¹⁶ bytes)
- Early Intel 80x86 chips addressed offsets of 'segments'
 - Thankfully segmented memory model has died off in x64
- Intel 80386 added **page-based memory mapping**



Page-based Memory Access



Challenges in Virtualizing Hardware Page Tables

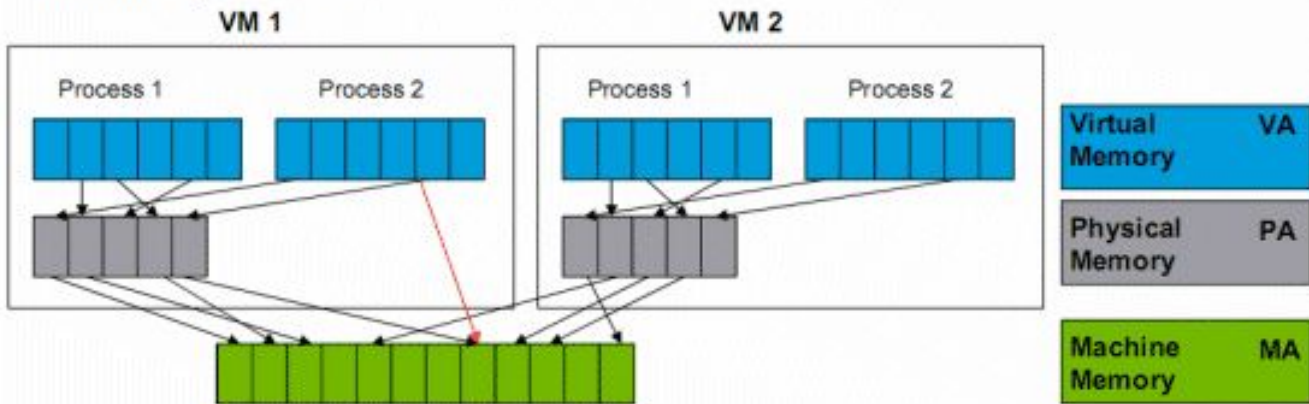
- Hypervisor has no chances to intercept on TLB misses
- Solutions
 - **Software-based:** Shadow Paging
 - **Hardware-based:** Hardware support for virtualizing memory



Virtualizing Paged Memory - Shadow Page Table

Virtualizing Virtual Memory


Shadow Page Tables



Virtualizing Paged Memory - Nested Paging

Hardware Support

Nested/Extended Page Tables

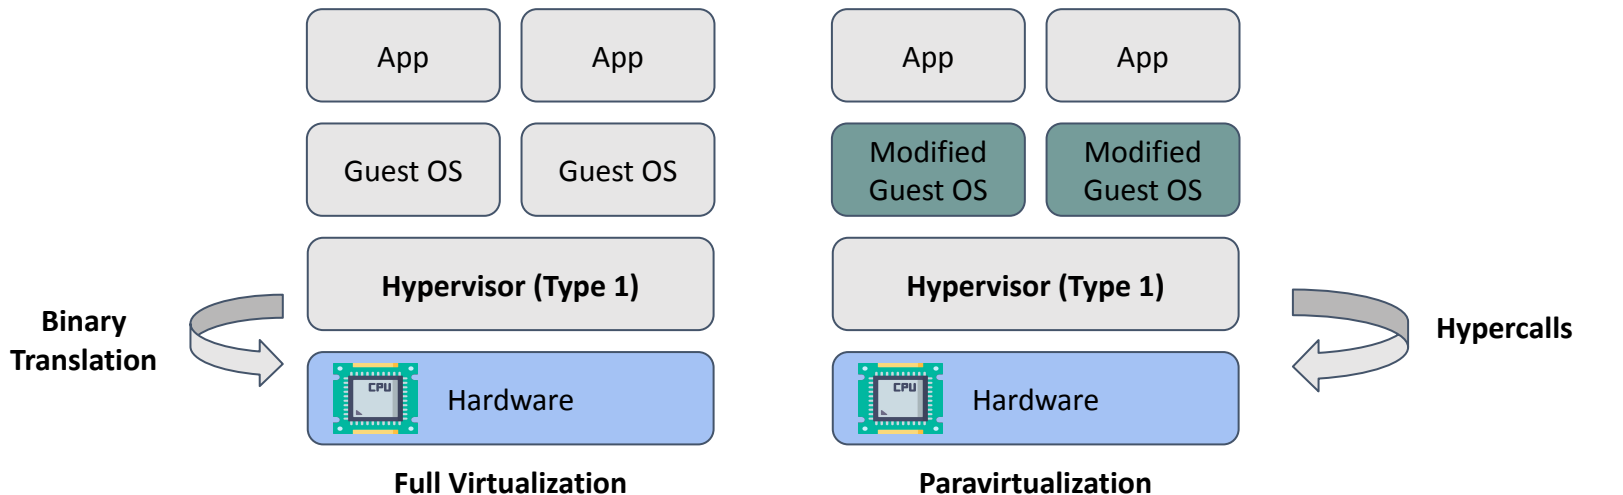


TLB Tagging

- Traditionally, every time a hypervisor switched between different VMs, the VM and its data structure had to be flushed out of the TLB
- Intel and AMD have facilitated TLB tagging since 2008
 - Intel Virtual Processor IDs (VPIDs) allow VMM to assign VM IDs
- Flush TLB entries of a particular VM only
 - So switching between VMs and VMM may leave TLB entries
 - Significant boost to memory access speed



Paravirtualization





Vagrant

- A tool for automatically creating and configuring VMs via simple scripts
- Operates on a **Vagrantfile** that contains instructions for how to download and run a particular VM image
- Requires a “provider” on the host machine to start and manage the VM
 - E.g., Virtualbox, QEMU, Hyper-V, ...



Agenda for Today

- OS Virtualization
 - No more hardware details!
 - Features enabling OS Virtualization
 - **Namespaces**
 - **Cgroups**
 - History of OS Virtualization
- Lot of Demos!
- Readings
 - Recommended: Demystifying Containers Part I
 - Optional:
 - <https://web.archive.org/web/20140906153815/http://www.cs.bell-labs.com/sys/doc/names.html>
 - <https://www.kernel.org/doc/html/latest/admin-guide/cgroup-v1/cgroups.html>
 - <https://www.redhat.com/sysadmin/pid-namespace>





Agenda for Today

- Quiz 1 is out!
 - 10 questions T/F, MC and multiple answers
- Presentation date signup
 - Please submit ASAP! FCFS!
 - <https://docs.google.com/spreadsheets/d/10x6tRuOoOK8TEUv3hZjCKSicSoRxcDEewkxTFiuka5w>






OS Virtualization

Types of Virtualization

- ~~Application Virtualization~~
 - E.g., JVM, PVM
- **OS Virtualization**
 - Containers
- ~~System Virtualization~~
 - Virtual Machines
- Network Virtualization
- Storage Virtualization
- Many more...



Different Types of Software Isolation

- Application: In-application isolation
 - Thread: Sharing memory within one process
 - Process: Each process has its own address space
 - **User space: All of the code in an OS that lives outside of the kernel**
 - Virtual Machine: What we have covered so far
-
- **OS virtualization is to isolate separate user spaces**
 - Assuming everything below and including the Kernel is shared



Why Isolate User space?

- Less resource intensive if the kernel is secure
 - (i.e., cheaper than needing to run VMs containing OS kernels)
- One example: Android
 - Allocates a unique user ID for every app
 - Each application's processes are (assumed to be) isolated
- **Another example: Containers (Docker)**



Two Features enabling OS Virtualization

- **Namespace**
 - Allows different view of the system for different process
- **cgroup**
 - Controls and isolates the resource usage (CPU, memory, disk I/O, network, etc.) for different processes
- Both are Linux Kernel features that are crucial for OS virtualization
- Warning: Lot of OS details coming up





Namespace

Namespace Overview

- Definition
 - Enable a process (or several processes) to have different views of the system than other processes
- Inspired by Plan 9 Distributed System from Bell Labs
 - Part of optional reading
- Originated in 2002 with “mnt” namespaces
- Now there are 8 kinds of namespaces



Kinds of Namespaces

- **pid: Process IDs**
- **mnt: Mount points, filesystems**
- **net: Virtualizing the network stack**
- **ipc: Interprocess communications**
- **uts: Unix Time Sharing for different hostname**
- **user: User identification**
- **cgroup: Identity of the control group (will cover this in later slides)**
- **time: Allow processes to see different times**



Namespace Demo

- To query namespaces: `ls -al /proc/<pid>/ns`
- By default, all “native” processes are placed in the same default namespaces
- How to enable separate namespaces?



Namespace Implementation

- Three ways: three system calls are used for namespaces:
- **clone()**: creates a new process and a new namespace; the process is attached to the new namespace.
 - Process creation and process termination methods, `fork()` and `exit()` methods, were patched to handle the new namespace `CLONE_NEW*` flags.
- **unshare()**: does not create a new process; creates a new namespace and attaches the current process to it.
 - `unshare()` was added in 2005, but not for namespaces only, but also for security. see “new system call, `unshare`” : <http://lwn.net/Articles/135266/>
- **setns()**: a new system call for joining an existing namespace



Digression: What are System Calls?

- A way for programs to interact with the operating system
- Provides the services of the operating system to the user programs via APIs
- Services Provided by System Calls :
 - Process creation and management (`fork()`, `exit()`, `wait()`)
 - Main memory management (`mmap()`)
 - File Access, Directory and File system management (`open()`, `read()`, `write()`, `close()`)
 - Device handling(I/O) (`ioctl()`, `read()`, `write()`)
 - Protection (`chmod()`, `chown()`)
 - Networking (`socket()`) etc.
- Types of System Calls : There are 5 different categories of system calls
 - Process control: end, abort, create, terminate, allocate and free memory.
 - File management: create, open, close, delete, read file etc.
 - Device management
 - Information maintenance
 - Communication



Namespace Implementation Details

- To support namespace, a member named `nsproxy` was added to the process descriptor called `struct task_struct`
 - `nsproxy` includes various namespaces: `uts_ns`, `ipc_ns`, `mnt_ns`, `pid_ns`, `net_ns`;
 - <https://elixir.bootlin.com/linux/latest/source/include/linux/nsproxy.h>
 - <https://elixir.bootlin.com/linux/latest/source/include/linux/sched.h>
- A method `task_nsproxy(struct task_struct *tsk)` used to access the `nsproxy` of a particular process
- There is an initial, default namespace for each namespace



Demo: PID Namespace

- Run the following commands
 - `unshare --fork -p /bin/bash`
 - `sleep 8000 &`
 - `ps -ef | grep sleep`
 - `kill <pid>`
 - `lsns |grep bash`
 - `sudo cat /proc/<pid>/status |grep NSpid`
- This create a new PID namespace by `unshare()` syscall and call `execvp()` for invoking `bash`
- Note that you can't kill other processes in different PIDs
- Building block for Container processes




Network Namespaces

- A network namespace is logically another copy of the network stack, with its own routes, firewall rules, and network devices.
- Network namespace includes all network stack information, like:
 - Loopback device
 - SNMP stats (netns_mib)
 - All network tables: routing, neighbors, etc
 - All sockets – /procfs and /sysfs entries
- Building block for Container networking




Demo: Network Namespaces

- Create two namespaces, `sfo` & `nyc`:
 - `ip netns add sfo`
 - `ip netns add nyc`
 - `ip netns list`
- Verify that two network namespaces are created:
 - `/var/run/netns/sfo`
 - `/var/run/netns/nyc`
- Which syscall is involved here?
 - `clone()`, `unshare()` or `setns()`?




Connecting Two Interfaces

- You can communicate between two network namespaces by:
 - Creating a pair of network devices (veth) and move one to another network namespace.
 - veth (Virtual Ethernet) is like a pipe




Demo: Connecting Two Interfaces

- Create a veth pair
 - `ip link add veth-sfo type veth peer name veth-nyc`
 - `ip link list | grep veth`
- Assign them to different network namespace:
 - `ip link set veth-sfo netns sfo`
 - `ip link set veth-nyc netns nyc`
- Run two processes associated with these two namespaces
 - `ip netns exec sfo ip link`
 - `ip netns exec nyc ip link`



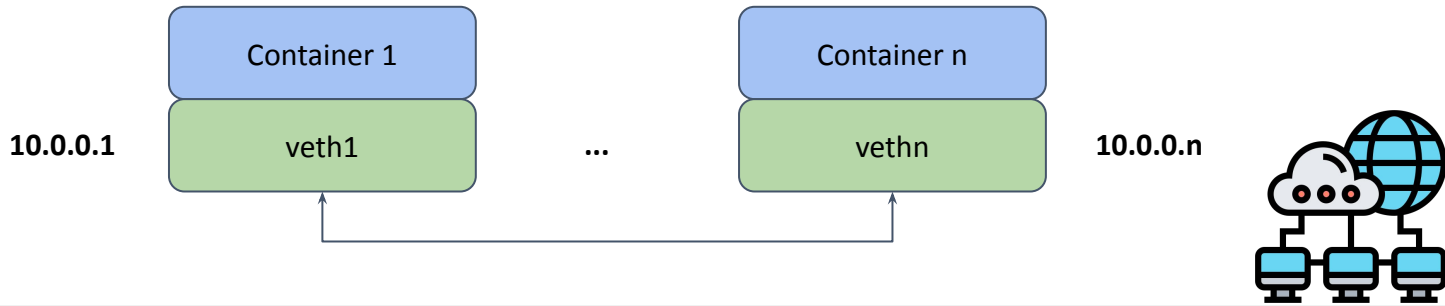
Demo: Connecting Two Interfaces

- Give the interfaces IP addresses
 - `ip netns exec sfo ip address add 10.0.0.1/24 dev veth-sfo`
 - `ip netns exec sfo ip link set veth-sfo up`
 - `ip netns exec nyc ip address add 10.0.0.2/24 dev veth-nyc`
 - `ip netns exec nyc ip link set veth-nyc up`
- Now check the IP addresses
 - `ip netns exec sfo ip addr`
 - `ip netns exec nyc ip addr`
- Now try to ping one another
 - `ip netns exec sfo ping 10.0.0.2`
 - `ip netns exec nyc ping 10.0.0.1`
- Clean up namespaces
 - `ip netns del nyc sfo`



Food for Thought: More Than 2 Containers?

- How to connect multiple containers?
- Our demo was for 1 to 1 network connection
- Hint: Use Linux Bridge




Mount (Filesystem) Namespaces

- Represented by `mnt_ns` (`mnt_namespace` object) in `nsproxy`
- For Linux, all previous mounts will be visible (inherited) in the new mount namespace, but...
 - mounts/unmounts in the new mount namespace are invisible to the rest of the system
- To create a new mount namespace
 - `unshare -m /bin/bash`
- How to specify a new root file system to a process
 - **chroot**: relink the root directory of the process to a new root directory (i.e., which includes a complete new file system of a container)




“Old School” Chroot Jail

- Unix servers have to handle users that may be malign
 - Common historical example was running public FTP servers
 - Anonymous users could log into those servers
 - FTP as a protocol allows quite a lot of power over the server
 - Needed to cut down what anonymous users could do
- Solution: change the perceived root directory of the filesystem
 - i.e., a ‘chroot jail’: changes the set of available executables
 - The process and its children can’t access the files outside the new root.
 - Unix accesses binaries from /bin, libraries from /lib, etc.
 - Changing the meaning of / mitigates many vulnerabilities



BSD Jails: OS-level virtualization since Y2K

- BSD Jail takes resource partitioning beyond the filesystem
 - Isolates process IDs, root user, network, device access
 - Also uses a chroot jail to effect filesystem isolation
- Helps avoiding privilege escalation
 - Successful break in to server can't scan filesystem for vulnerabilities, e.g., reading /etc/shadow and trying to crack weak passwords
- Many operations are blocked within BSD Jails
 - E.g., loading kernel modules, changing network interfaces, mounting and unmounting filesystems, etc.





Cgroup (Control Group)

Linux Cgroup

- This work was started by engineers at Google in 2006 under the name "process containers; in 2007, renamed to "Control Groups".
- Defines parameters about the resource use of a set of processes, e.g.:
 - Limit total memory available to group of processes
 - Indicate non-even share of device input/output priority
 - Affect CPU scheduling to the group
 - cgroups also can assist accounting for resource use
 - cgroups can be applied hierarchically
- cgroups can facilitate starting / stopping processes
 - Important for snapshot functionality



Cgroup Subsystems

- Kernel modules that are used to control the access that cgroups have to various system resources
- Example 10 cgroup subsystems for Redhat
 - `blkio` — sets limits on input/output access to and from block devices such as physical drives
 - `cpu` — uses the scheduler to provide cgroup tasks access to the CPU.
 - `cpuacct` — generates automatic reports on CPU resources used by tasks in a cgroup.
 - `cpuset` — assigns individual CPUs (on a multicore system) and memory nodes to tasks in a cgroup.
 - `devices` — allows or denies access to devices by tasks in a cgroup.
 - `freezer` — suspends or resumes tasks in a cgroup.
 - `memory` — sets limits on memory use by tasks in a cgroup and generates reports on memory usage.
 - `net_cls` — tags network packets with a class identifier (classid) that allows the Linux traffic controller (`tc`) to identify packets originating from a particular cgroup task.
 - `net_prio` — provides a way to dynamically set the priority of network traffic per network interface.
 - `ns` — the *namespace* subsystem.
 - `perf_event` — identifies cgroup membership of tasks and can be used for performance analysis



Cgroup Virtual File System (VFS)

- Cgroups uses a Virtual File System
- All entries created in it are not persistent and deleted after reboot.
- All cgroups actions are performed via file system actions (create/remove directory, reading/writing to files in it, mounting/mount options).



Mounting Cgroups

- In order to use a filesystem, it must be mounted
- A control group can be mounted anywhere on the filesystem. (e.g., Systemd uses `/sys/fs/cgroup`.)
- When mounting, we can specify with mount options (`-o`) which subsystems we want to use
 - `mkdir /cgroup/memtest`
 - `mount -t cgroup -o memory test /cgroup/memtest/`



Mounting Cgroups

- Under each new cgroup which is created, some common files are always created
 - `tasks`: list of pids which are attached to this group.
 - `cgroup.procs`: list of thread group IDs (listed by TGID) attached to this group.
- Each subsystem adds specific control files for its own needs
 - `memory.max_usage_in_bytes`
 - `memory.limit_in_bytes`
 - `memory.kmem.tcp.limit_in_bytes`
 - `memory.kmem.tcp.max_usage_in_bytes`
 - ...



Cgroup Demo: cpuset

- `cpusets` provide a mechanism for assigning a set of CPUs and Memory Nodes to a set of tasks
- Creating a cpuset group is done with:
 - `mkdir /sys/fs/cgroup/cpuset/group1`
 - `echo 0 > /sys/fs/cgroup/cpuset/group1/cpuset.cpus`
 - `echo 0 > /sys/fs/cgroup/cpuset/group1/cpuset.mems`
 - `echo #pid > /sys/fs/cgroup/cpuset/group1/tasks`



Cgroup Demo: Memory

- `mkdir /sys/fs/cgroup/memory/group1`
- `echo #pid > /sys/fs/cgroup/memory/group1/tasks`
- `echo 10M > /sys/fs/cgroup/memory/group1/memory.limit_in_bytes`
- What would happen if you run a process demanding more than 10 M memory?
- Optional Reading
 - https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux/6/html/resource_management_guide/sec-memory



Summary

- Namespaces and Cgroups are building blocks for containers
- The container runtime is responsible for
 - Creating namespaces
 - Creating cgroups
 - Executing a process within the given namespace and cgroups, which makes it a container!



Agenda for Today

- OS Virtualization
 - No more hardware details!
 - Features enabling OS Virtualization
 - **Namespaces**
 - **Cgroups**
 - History of OS Virtualization
- Lot of Demos!
- Readings
 - Recommended: Demystifying Containers Part I
 - Optional:
 - <https://web.archive.org/web/20140906153815/http://www.cs.bell-labs.com/sys/doc/names.html>
 - <https://www.kernel.org/doc/html/latest/admin-guide/cgroup-v1/cgroups.html>
 - <https://www.redhat.com/sysadmin/pid-namespace>





TODOs!

- HW 1 Will be released next class
- Project
- Quiz 1





Questions?

