



# COEN 241

# Introduction to Cloud Computing

## Lecture 8 - Kubernetes





# Lecture 7 Recap

- Kata Container
- Orchestration in Cloud
- Infrastructure as Code
- CoreOS
- Readings
  - Recommended: None
  - Optional:
    - <https://www.youtube.com/watch?v=4gmLXyMeYWI>
    - <https://www.stackhpc.com/kata-io-1.html>
    - <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=9198653>
    - <https://www.techrepublic.com/article/simplifying-the-mystery-when-to-use-docker-docker-compose-and-kubernetes/>

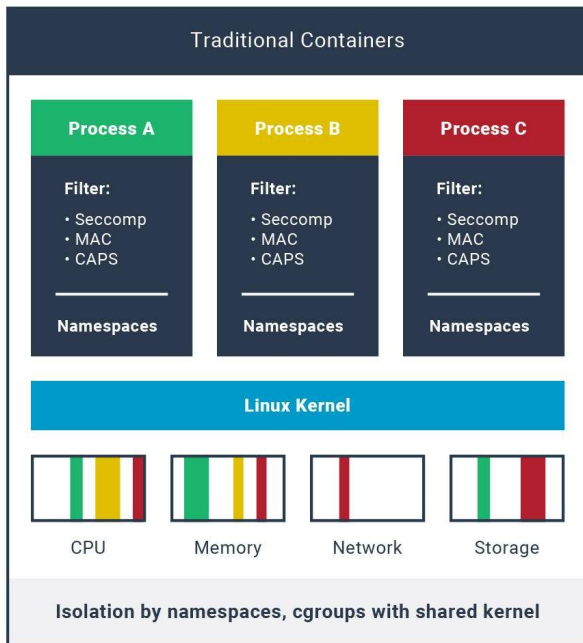
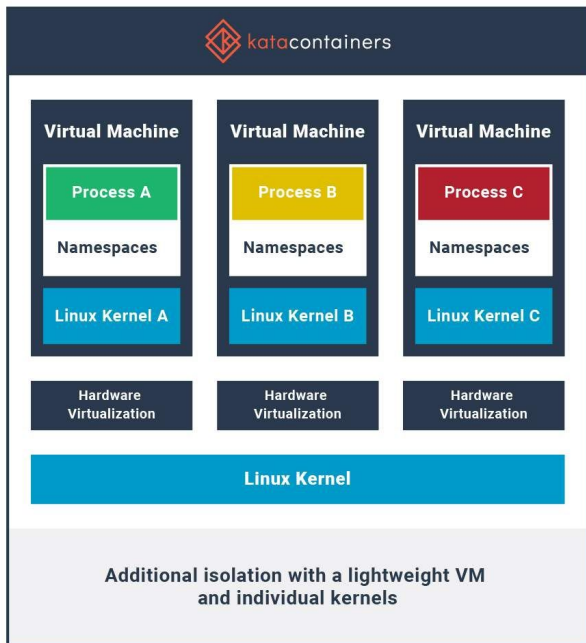


# What is Kata Container?

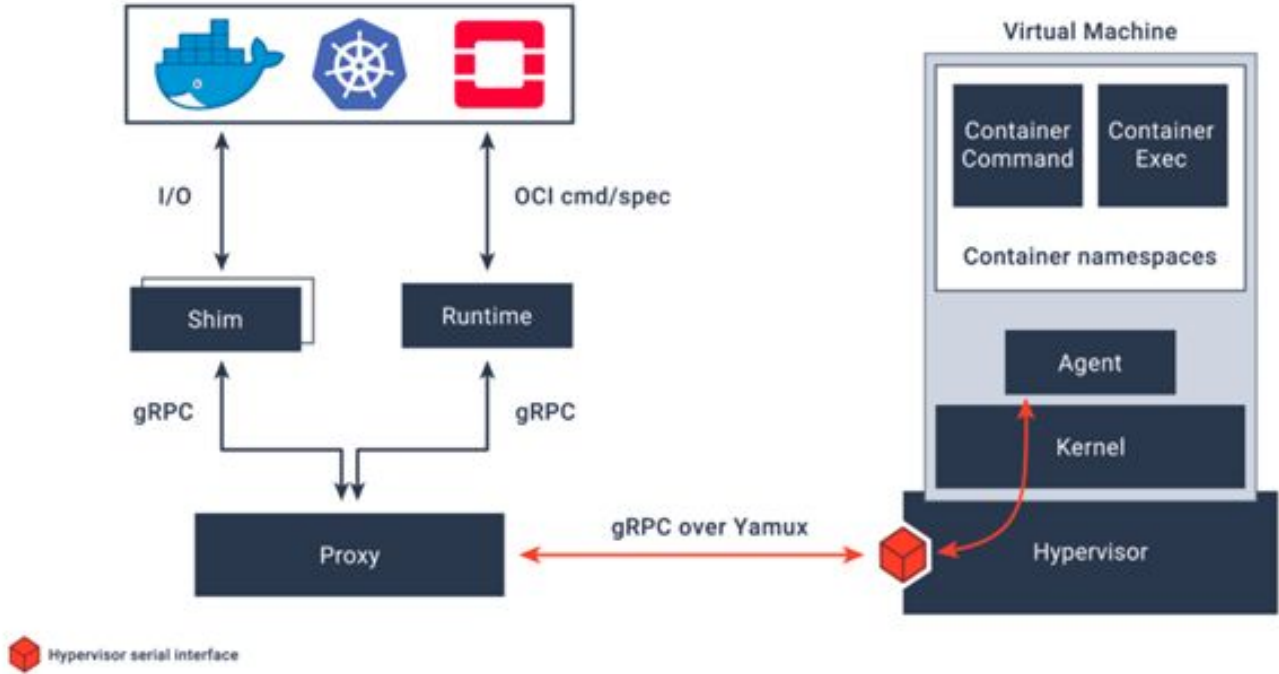
- Released in 2017
  - From the merge of Intel's Clear Containers and Hyper's runV
  - "Wraps" containers into dedicated virtual machines
  - OCI runtime implementation can be plugged into the container engine
    - Docker
  - Can consume existing container images
- Kata is a container runtime
  - Can still be coupled with other Docker platforms



# Kata Architecture



# Kata Workflow



# Orchestration

- Definition: Automatic management of computer systems
  - Deployment and configuration
  - Interconnection and coordination
  - Monitoring
    - Can also configure to hook up monitoring for management
    - E.g., Deploy more resources as load increases
- Growing set of very good (open-source) solutions:
  - **Machine focused:** e.g., Puppet, Terraform, Ansible, Salt, Chef...
  - **Cloud-based:** e.g., AWS CloudFormation, Terraform
  - **Container clusters:** e.g., Kubernetes



# Automation vs Orchestration

- **Automation:** completing a **single task** or function without human intervention
- **Orchestration:** Managing a large-scale virtual environment or network by orchestrating the scheduling and integration of automated tasks between complex distributed systems and services
  - Simplifies interconnected workloads, repeatable processes, and operations.
- To simplify, **Automation** refers to a single task vs **Orchestration** arranges multiple tasks to optimize a workflow



# Declarative Configuration Management

- Declarative tools specify the desired target state
  - E.g., Can I have a coffee on my desk at 9AM on Monday morning?
- The means to reach target state is up to the configurations
  - Can take corrective action to react to drift in machine's state
- State specification will be a domain specific language (DSL)
- Some example FOSS systems with large user communities
  - Puppet, Terraform, SaltStack





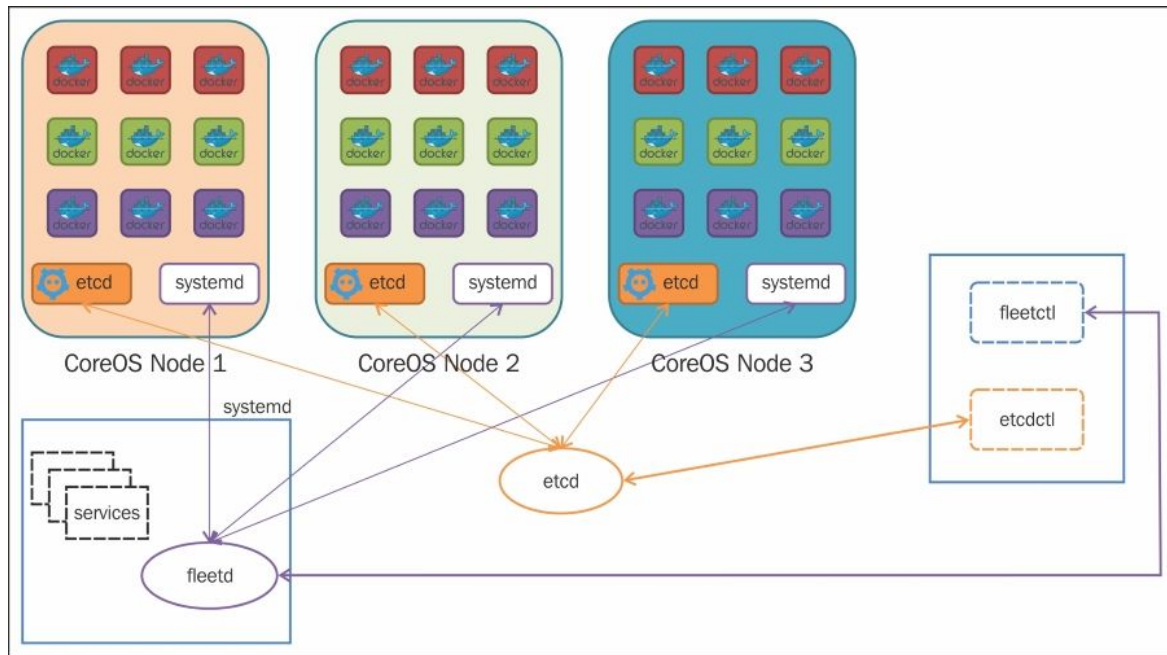


# Imperative Configuration Management

- Also called 'procedural', i.e., specifying steps to run:
  - Usually written in chunks of code in configuration system authors' favorite PL
- Some example systems:
  - Ansible (Py), Chef (Ruby), Saltstack
- Can write imperative code to have declarative effect



# Core OS Architecture



# Agenda for Today

- Kubernetes
  - Concepts
  - Architecture
- Minikube
- Readings
  - Recommended: None
  - Optional:
    - <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
    - <https://labs.play-with-k8s.com/>
    - <https://minikube.sigs.k8s.io/docs/handbook/>
    - <https://circleci.com/blog/what-is-yaml-a-beginner-s-guide/>





# Kubernetes Overview

# What is Kubernetes?

- An open-source container-orchestration system for automating computer application deployment, scaling, and management.
- It is a Greek word meaning helmsman or pilot.
- Often called K8s, which is derived by replacing the eight letters of “ubernet” with the digit 8.
- Project emerged from Google in 2014; v1.0 released 2015
  - Related to Borg—Google’s (secret) internal container scheduler
  - K8s is implemented in Go, in contrast to Borg’s C++



# Why is Kubernetes Useful?

- Provides an easy way to scale your containers
- Avoids cloud vendor lock-in
- Allows easy multi-cloud adoption
- Better management of application
- Works with many different container technologies
- It is a BIG piece of software with lot of moving parts
  - We will give a short, high-level overview of it
  - There are numerous tutorials and classes dedicated to Kubernetes





# Interacting with Kubernetes

- kubectl
- APIs



## Some Kubernetes Terms

- **Pods:** tightly coupled **set of containers**; smallest unit of scheduling
- **Node:** A Pod always runs on a Node. A Node is **a worker machine** in Kubernetes and may be either **a virtual or a physical machine**.
- **Services:** An abstraction which defines a logical set of Pods and a policy by which to access them
  - Sometimes this pattern is called a micro-service architecture
- **Volumes:** Persistent storage that can be shared between containers
- **Namespaces:** Multiple virtual cluster in a same physical cluster





# Kubernetes Pods

- Pods are the smallest deployable units of computing that you can create and manage in Kubernetes
- Group of one or more containers with **shared storage** and **network resources**, and a specification for how to run the containers
  - A group of Docker containers with shared namespaces and file system volumes
- A Pod's contents are always co-located and co-scheduled, and run in a shared context




# Kubernetes Pods

- Pods are assigned an IP address, for networking
  - All containers within the pod share that address and its ports
- Pods provide app. storage (volumes) to containers
- Pods usually created by controllers, and not directly
- List and show detailed information of the Pods by running
  - `kubectl get pods`
  - `kubectl describe pods`



# Kubernetes Pod Overview






## Example Pod Definition (PodSpec)

```
apiVersion: v1
kind: Pod
metadata:
  name: multi-pod
  label: myApp
spec:
  volumes:
    - name: html
      emptyDir: {}
  containers:
    - name: 1st
      image: nginx
      volumeMounts:
        - name: html
          mountPath: /usr/share/nginx/html
    - name: 2nd
      image: debian
      volumeMounts:
        - name: html
          mountPath: /html
      command: ["/bin/sh", "-c"]
      args:
        - while true; do
            date >> /html/index.html;
            sleep 1;
          Done
```

# Kubernetes Nodes

- A set of worker machines that run containerized applications
- Every cluster has at least one worker node
- When you deploy K8s, you get a **cluster**
- `kubectl get nodes`



# Kubernetes Services

- A **logical abstraction** for a deployed group of pods in a cluster
  - which all perform the same function
- Kubernetes way of configuring a proxy to forward traffic to a set of pods
- Since pods are ephemeral, a service enables a group of pods to be assigned a name and unique IP address (clusterIP)
- As long as the service is running on that IP address, it will not change.
- Services also define policies for the access to Pods



# Types of Kubernetes Services

- **ClusterIP** (default): Exposes a service which is only accessible from within the cluster.
- **NodePort**: Exposes a service via a static port on each node's IP.
- **LoadBalancer**: Exposes the service via the cloud provider's load balancer.
- **ExternalName**: Maps a service to a predefined externalName field by returning a value for the CNAME record.





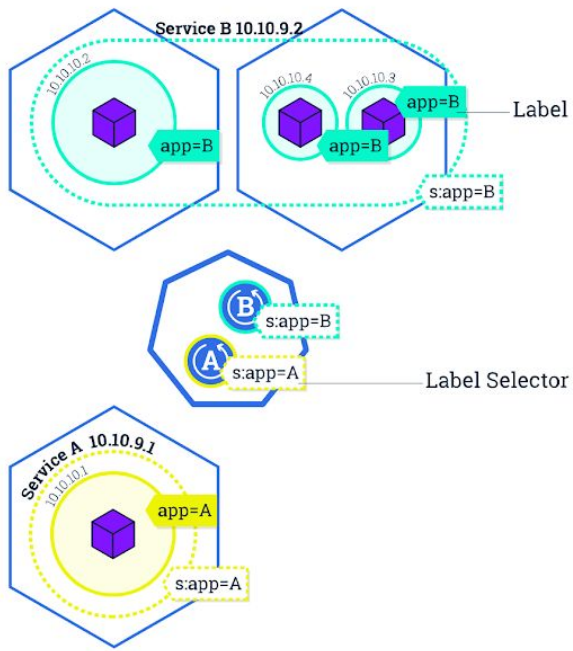
## Example Service Definition

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: ClusterIP
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```





# Service Overview



# Stateless vs Stateful Applications

- Stateless applications scale easily: just start more pods
  - E.g., web servers presenting read-only workloads
- Stateful apps are more difficult, e.g.:
  - Databases having primary and secondary instances
  - Machine Learning Applications
- Kubernetes StatefulSet controllers enables stateful applications
  - Provides durable storage
  - Persistent ID for each Pod
  - <https://kubernetes.io/docs/tutorials/stateful-application/basic-stateful-set/>

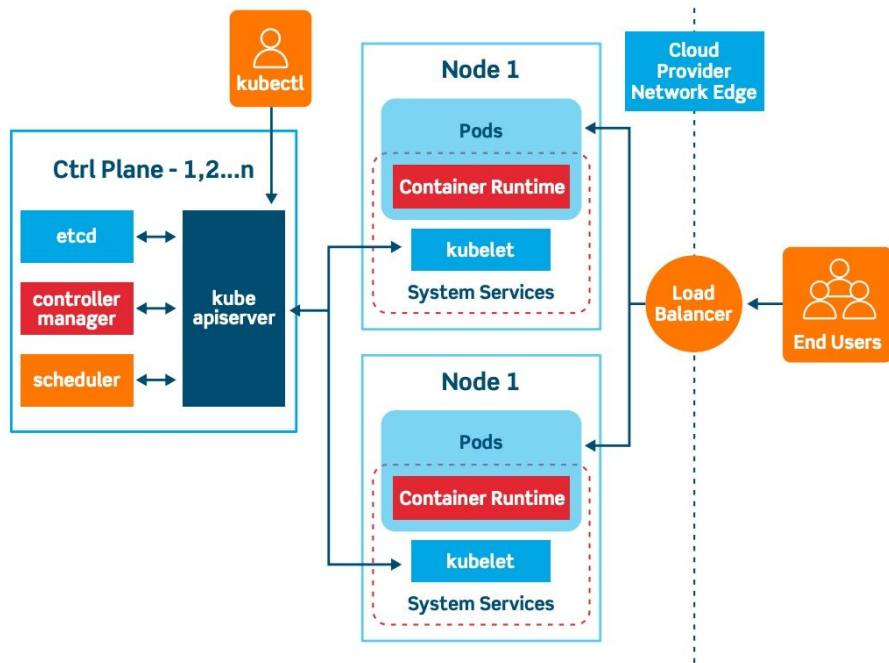




# Kubernetes Architecture

# Kubernetes Architecture Overview

- Kubernetes cluster consists of:
- Master Node(s) (control plane)
  - With a distributed storage (etcd)
- A number of cluster nodes



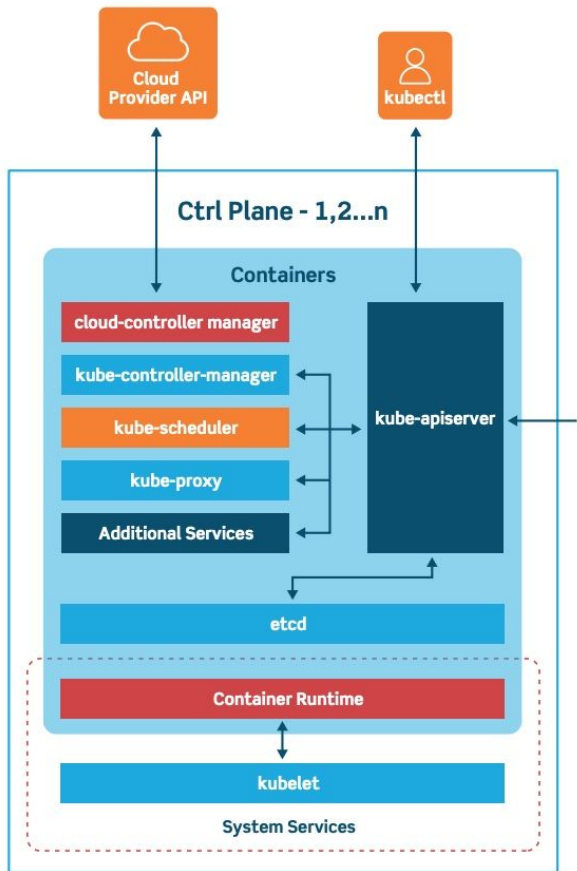
# Kubernetes Master Node

- A logically centralized control point
- Maintains a record of all Kubernetes objects
  - Manages object states, responding to changes in the cluster
  - Works to make the actual state of system objects match the desired state
- Consists of the following components
  - **API server:** Allows Kubernetes cluster to be controlled from a user
  - **Controller manager:** Checks replication, status of nodes and state of cluster
  - **Scheduler:** Allocates pods waiting to run nodes
  - **etcd:** Consistent repository of configuration information



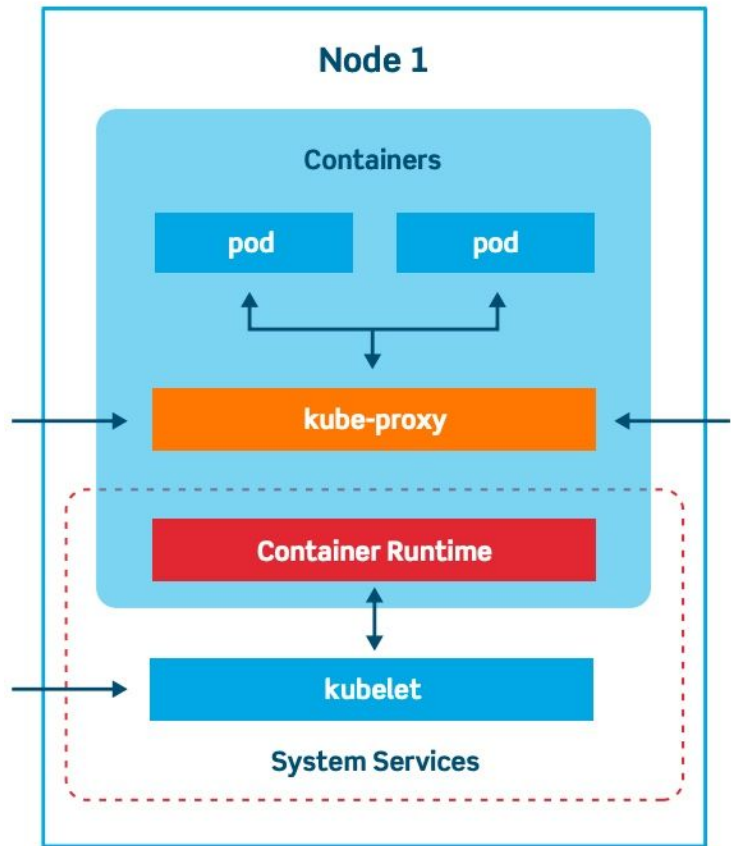
# Kubernetes Master Node

- Can be one or more master nodes
- Can accept APIs from cloud providers
- Also provides proxy for network access
- **etcd** can be configured for multi-node distributed configuration manager



# Kubernetes Cluster Node

- Nodes can run pods but also,
  - kube-proxy: provides network services; leveraging OS facilities
  - **kubelet**: checks on health of containers within a pod
  - cAdvisor: provides statistics about container resource use
- **Kubelet** is the primary and most important controller
  - Drives the container execution layer



# Kubelet

- The kubelet is the primary "node agent" that runs on each node
- It can register the node with the kube apiserver using one of:
  - The hostname
  - A flag to override the hostname
  - Specific logic for a cloud provider
- The kubelet ingests a PodSpec and runs each pod







# Extension/Alternatives to Kubernetes

# K8s as a Service

- K8s can manage your containers, but how to set it up given a IaaS?
  - Need VMs running master and the nodes
- When using a single cloud provider it can provide a range of options
  - For example, AWS offer a range of options such as
    - AWS Fargate provides a complete container service
    - AWS EKS provides control plane; you set up K8s nodes on EC2
    - Use EC2 to deploy all the components if you want full control

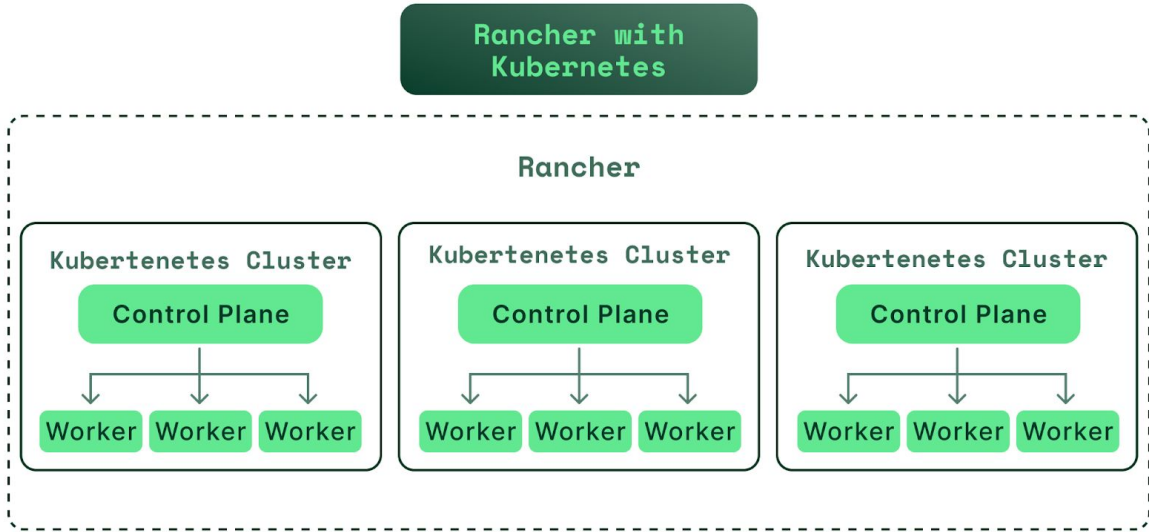


# K8s in Multi-Cloud Application

- For multi-cloud applications, you need a different service
- Tools like Rancher provide for **multi-K8s** cluster management
  - Can even import existing cluster!
  - <https://rancher.com/>
  - <https://www.youtube.com/watch?v=LX0zVh-AYm4>
- This works because different cloud providers' container services are very similar!



# Rancher versus Kubernetes



# Rancher versus Kubernetes

- **Kubernetes Features**
  - Kubernetes based platform is easily migratable across cloud providers
  - Easier to scale as compared to traditional applications hosted in virtual machines (VM)
  - Easy to control cluster density and autoscaling
  - In node failures, pods are automatically rescheduled to other nodes
  - Can be hard to manage multiple clusters
- **Rancher Features**
  - Easy visualization to manage multiple clusters
  - Can easily create new clusters or add existing ones to it
  - Add a concept of projects for better grouping of namespaces
  - Permissions can be configured per project across clusters



## Terraform versus Rancher, Kubernetes, etc.

- Rancher can help deploy Kubernetes over bare metal
  - Also includes monitoring & security management tools
- However, you may need specific type of infrastructure nodes
  - E.g., configuring a GPU node on Amazon for deep learning
  - Containers can use this hardware, but what about container managers?
- Terraform is a stack below tools like Rancher, it:
  - Allows previous of the plans using a form of IAC
  - Can easily provision at level of particular instance type
  - Can work with Rancher and Kubernetes



# Why it's so hard to pick a “winner”?

- All are churning rapidly in what's provided, e.g.:
  - Rancher's original functionality replaced by Docker Swarm
  - CoreOS Linux's original fleet functionality replaced by K8s
- Then what to use? Consider your and **the developers' time**
  - Look to see whether a new tool can optimise your processes
  - when taking into account the cost of transition
- Must have **IAC** and continuous integration pipelines
  - IAC is a must!



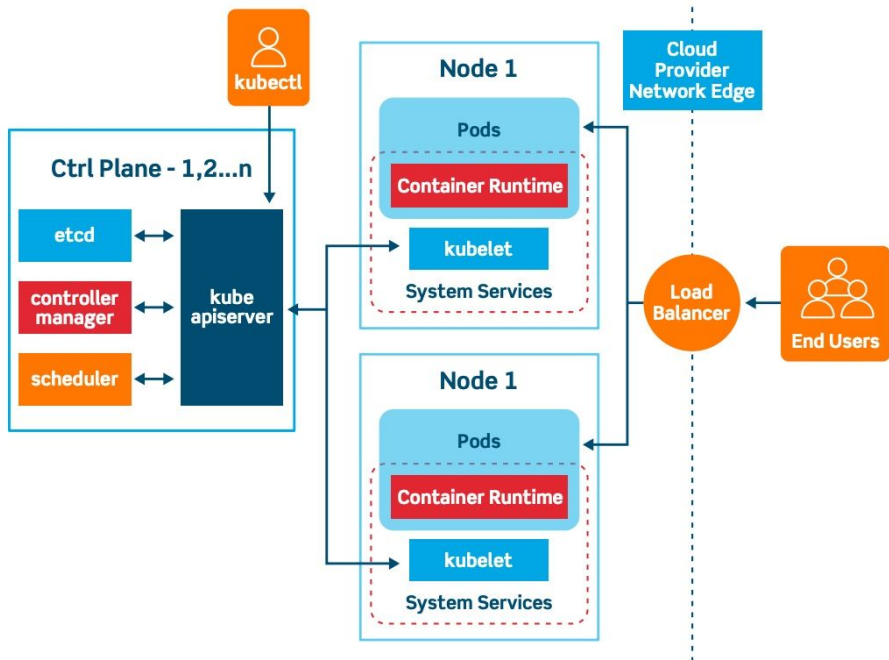


# Minikube



# Production Kubernetes Cluster

- Requires
  - Multiple Master Nodes
  - Multiple Worker Nodes
- Hard to try Kubernetes locally!
- Minikube to the rescue



# Minikube: A local Kubernetes setup

- <https://minikube.sigs.k8s.io/docs/start/>
- Running both master and worker processes **on the same node**
- Minimum requirement: 2 CPUs, 2GB of RAM, 20GB of free disk space
- Use **kubectl** to interact with the Minikube cluster
- **Kubectl**: A CLI for Kubernetes cluster to talk to the master
  - Not only Minikube, but any type of Kubernetes cluster
  - <https://kubernetes.io/docs/reference/kubectl/overview/>





# TODOs!

- Kubernetes Quiz!
  - <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
  - Up to Module 4 Only!
- HW 2
- Study for Midterms!





# Agenda for Today

- Kubernetes
  - Concepts
  - Architecture
- Minikube
- Readings
  - Recommended: None
  - Optional:
    - <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
    - <https://labs.play-with-k8s.com/>
    - <https://minikube.sigs.k8s.io/docs/handbook/>
    - <https://circleci.com/blog/what-is-yaml-a-beginner-s-guide/>





# Questions?

