



# COEN 241

# Introduction to Cloud Computing

## Lecture 9 - FaaS/Serverless





# Lecture 8 Recap

- Kubernetes
  - Concepts
  - Architecture
- Minikube
- Readings
  - Recommended: None
  - Optional:
    - <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
    - <https://labs.play-with-k8s.com/>
    - <https://minikube.sigs.k8s.io/docs/handbook/>
    - <https://circleci.com/blog/what-is-yaml-a-beginner-s-guide/>



# What is Kubernetes?

- An open-source container-orchestration system for automating computer application deployment, scaling, and management.
- It is a Greek word meaning helmsman or pilot.
- Often called K8s, which is derived by replacing the eight letters of “ubernete” with the digit 8.
- Project emerged from Google in 2014; v1.0 released 2015
  - Related to Borg—Google’s (secret) internal container scheduler
  - K8s is implemented in Go, in contrast to Borg’s C++



# Why is Kubernetes Useful?

- Provides an easy way to scale your containers
- Avoids cloud vendor lock-in
- Allows easy multi-cloud adoption
- Better management of application
- Works with many different container technologies
- It is a BIG piece of software with lot of moving parts
  - We will give a high-level overview of it



## Some Kubernetes Terms

- **Pods:** tightly coupled set of containers; smallest unit of scheduling
- **Node:** A Pod always runs on a Node. A Node is a worker machine in Kubernetes and may be either a virtual or a physical machine,
- **Services:** An abstraction which defines a logical set of Pods and a policy by which to access them
  - Sometimes this pattern is called a micro-service
- **Volumes:** persistent storage; can share between containers
- **Namespaces:** multiple virtual cluster in a same physical cluster






# Kubernetes Pods

- Pods are the smallest deployable units of computing that you can create and manage in Kubernetes
- Group of one or more containers with shared storage and network resources, and a specification for how to run the containers
  - A group of Docker containers with shared namespaces and file system volumes
- A Pod's contents are always co-located and co-scheduled, and run in a shared context



# Kubernetes Nodes

- A set of worker machines that run containerized applications
- Every cluster has at least one worker node
- When you deploy K8s, you get a **cluster**
- `kubectl get nodes`



# Kubernetes Services

- A **logical abstraction** for a deployed group of pods in a cluster
  - which all perform the same function
- Kubernetes way of configuring a proxy to forward traffic to a set of pods
- Since pods are ephemeral, a service enables a group of pods to be assigned a name and unique IP address (clusterIP)
- As long as the service is running on that IP address, it will not change.
- Services also define policies for the access to Pods





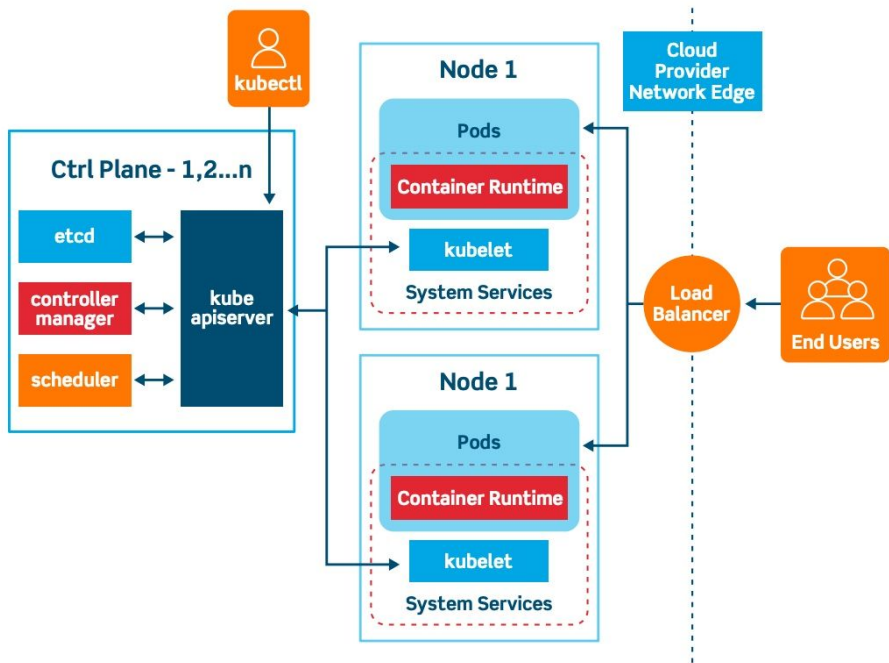
# Stateless vs Stateful Applications

- Stateless applications scale easily: just start more pods
  - E.g., web servers presenting read-only workloads
- Stateful apps are more difficult, e.g.:
  - Databases having primary and secondary instances
  - Machine Learning Applications
- Kubernetes StatefulSet controllers enables stateful applications
  - Provides durable storage
  - Persistent ID for each Pod
  - <https://kubernetes.io/docs/tutorials/stateful-application/basic-stateful-set/>



# Kubernetes Architecture Overview

- Kubernetes cluster consists of:
- Master Node(s) (control plane)
  - With a distributed storage (etcd)
- A number of cluster nodes



# Minikube: A local/test setup

- <https://minikube.sigs.k8s.io/docs/start/>
- Running both master and worker processes on the same node with a minimum requirement:
  - 2 CPUs or more, 2GB of free memory, 20GB of free disk space
- Use **kubectl** to interact with the Minikube cluster
- **Kubectl**: A CLI for Kubernetes cluster to talk to the master
  - Not only Minikube, but any type of Kubernetes cluster
  - <https://kubernetes.io/docs/reference/kubectl/overview/>





# Agenda for Today

- Function as a Service
- OpenFaaS
- Readings
  - Recommended: CCSA Chapter 4
  - Optional:
    - <https://aws.amazon.com/solutions/case-studies/netflix-and-aws-lambda/>





# Function as a Service (FaaS)

# What is Function as a Service?

- A cloud computing platform that allows customers to run and manage applications **without the complexity of building and maintaining the infrastructure** needed to run the application.
- Started in 2010 by a start up called PiCloud
  - <https://techcrunch.com/2010/07/19/picloud-launches-serverless-computing-platform-to-the-public/>
  - Acquired by Dropbox in 2013
- AWS Lambda is the first public cloud to offer FaaS in 2014
  - Followed by Google Cloud Functions, Microsoft Azure Functions, IBM/Apache's OpenWhisk in 2016 and Oracle Cloud Fn in 2017



# What is Function as a Service?

- Developers write independent code scripts known as “functions” and upload those functions to the FaaS
  - That’s why we call it “Function”-as-a-Service
- Once uploaded, the functions can be then triggered by some event or run on a schedule





# FaaS Example (in Python)

```
import os
import json

def lambda_handler(event, context):
    return {
        "statusCode": 200,
        "headers": {
            "Content-Type": "application/json"
        },
        "body": json.dumps({
            "message": "Hellow World"
        })
    }
```





# Serverless vs FaaS

- Serverless is a **cloud computing execution model** that abstracts away much of the complexity associated with infrastructure management from developers
- FaaS is a common form of serverless computing
  - i.e., Trend away from needing to considering servers at all
- FaaS is a platform providing serverless architecture deployment, orchestration and management

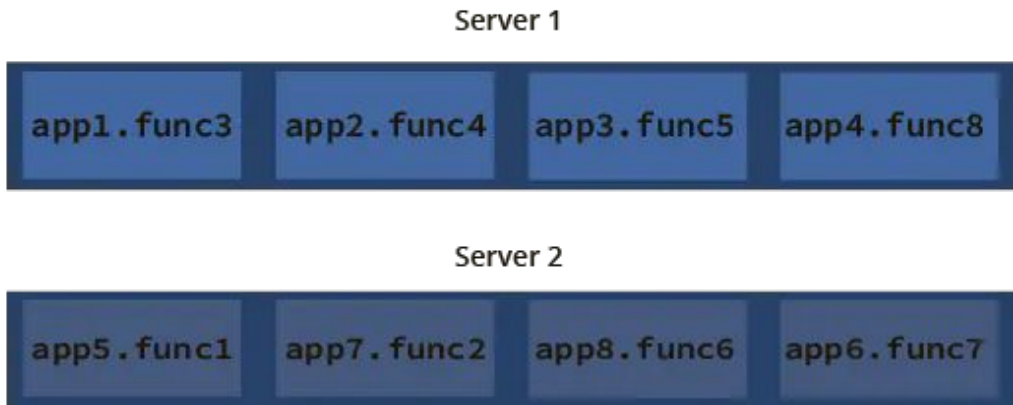


# Function as a Service vs Others

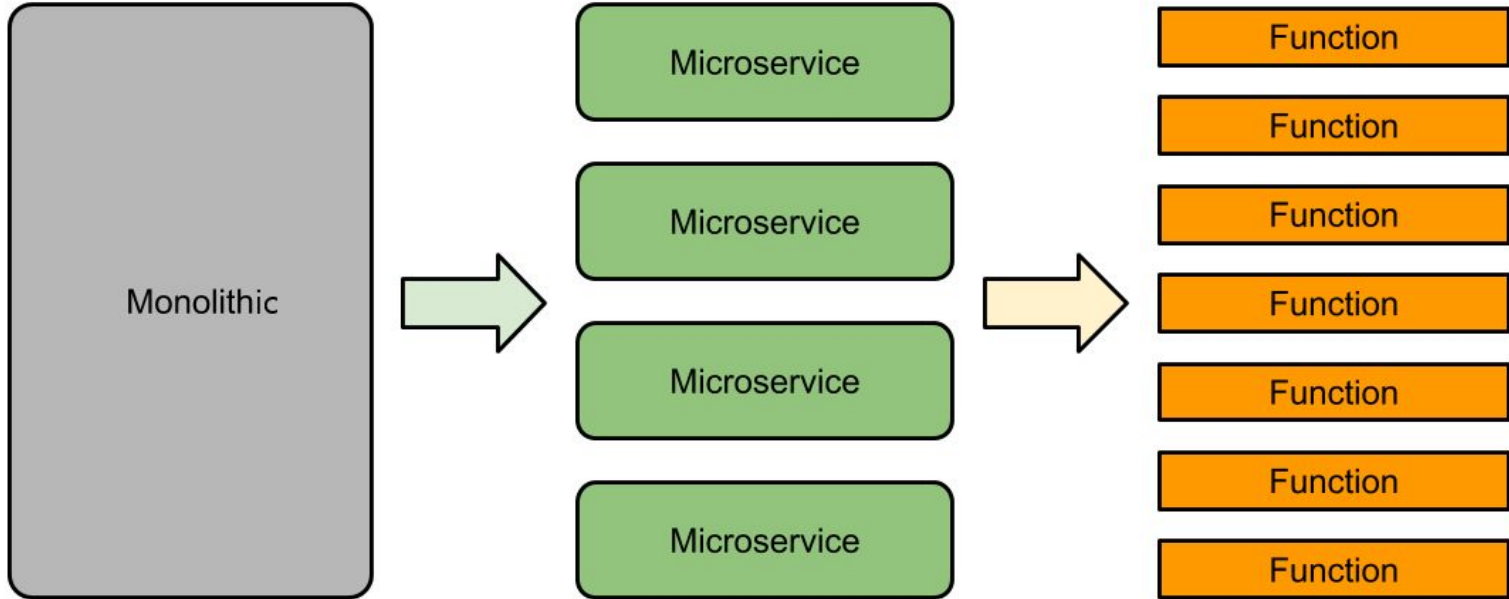
## Traditional architecture



## Serverless architecture



# Function as a Service vs Others



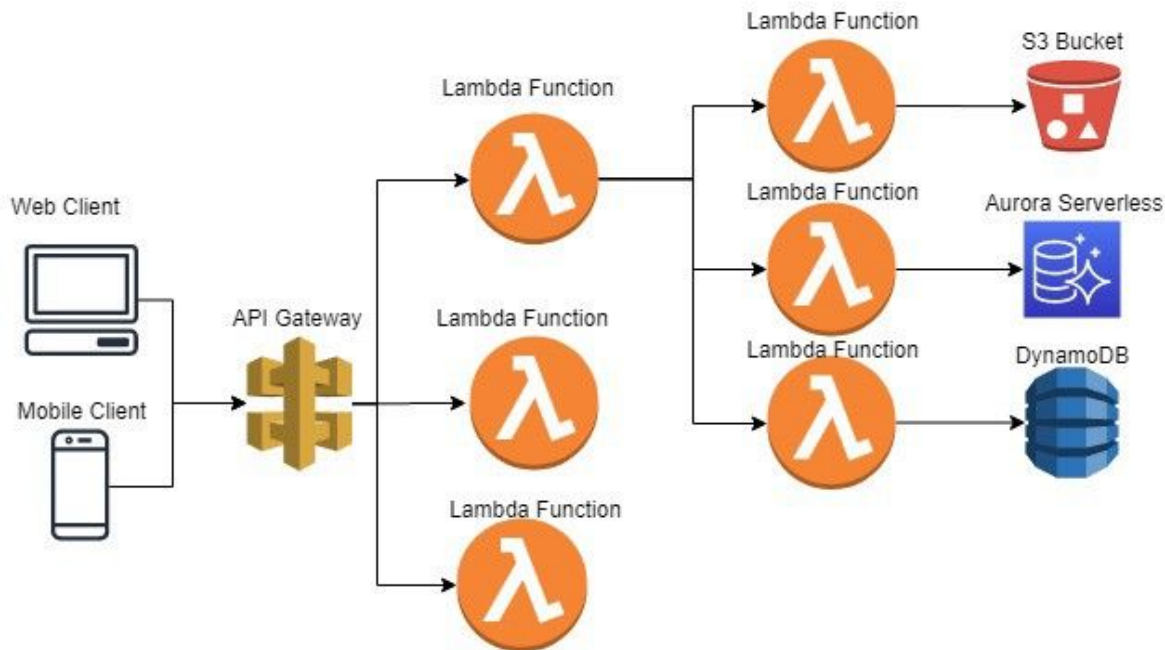


# Function as a Service vs Others

Private Cloud	IaaS Infrastructure as a Service	PaaS Platform as a Service	FaaS Function as a Service	SaaS Software as a Service
Function	Function	Function	Function	Function
Application	Application	Application	Application	Application
Runtime	Runtime	Runtime	Runtime	Runtime
Operating System	Operating System	Operating System	Operating System	Operating System
Virtualization	Virtualization	Virtualization	Virtualization	Virtualization
Server	Server	Server	Server	Server
Storage	Storage	Storage	Storage	Storage
Networking	Networking	Networking	Networking	Networking

Managed by the customer ■  
Managed by the provider ■

# Function as a Service Architecture



# Function as a Service Architecture

- Functions in FaaS are “stateless”, which means they cannot save and share session states between two different runs of the same function or different functions
  - Must use an external means of storing states
- Function can call other serverless functions as part of an application.
  - Need to communicate with other functions using asynchronous messaging




# FaaS Use Cases

- Many use cases
  - Web apps / Backends
  - Chatbots
  - Scheduled tasks
  - IT Automation
  - **Distributed stream processing systems (DSPS)**
  - **Event-driven programming**
  - **Internet of Things (IoT)**
  - **Edge Computing**
- Involving **short, self-contained** tasks with low memory requirement



# Distributed Stream Processing Systems (DSPS)

- Traditional DBs vs DSPSs
  - Databases: Runs queries when instructed to do so
  - DSPS: Computing is triggered by new tuples appearing
- Stream processing systems define a data flow graph
  - Data Source -> Operators -> Sinks
  - Operators: nodes that transform input streams to outputs
- Many high-quality, “scale-out”, open source DSPSs:
  - Apache Storm, Apache Spark Streaming, Apache Flink, ...
  - Scale-out: Adding more resources
  - Scale-up: Adding power to given resource





# Event-Driven Programming

- In IaaS, VM is yours, so your code is always running
- In PaaS, still usually a sense that your code is active
  - PaaS auto-scales the server instances that run your code
- In FaaS, your code operates in a **reactive style**
  - Reactive programming typically relies on callbacks
  - Some sort of shared event dispatcher issues callbacks
  - FaaS may have wide variance in function execution latency
    - Why?



# FaaS and Internet of Things

- IoT embodies ambition of all devices being networked
  - Devices including cars, streetlights, toasters, wireless sensors, etc.
- For sensor networks, want to offload data processing
  - To extend lifetime of battery-powered devices
  - Sensor data will be disseminated periodically, often at random intervals
- FaaS facilitates data transformations
  - Provides reliable endpoints for IoT devices to interact with on demand



# FaaS and Edge Computing

- Edge computing sits in between cloud and IoT
  - Often consists of full-size computing devices, widely distributed
  - ... but not at cluster scale (so not scale-out edge computing)
- Example: Content Delivery Networks (CDN)
  - Geographically distributed servers that speed up the delivery of web content by bringing it closer to where users are
- CloudFront CDN runs AWS Lambda functions
  - E.g., personalize web content within any AWS Region's DC
- Assist in short tasks without traveling all the way to the DCs



# Advantages of Function as a Service

- Increased developer productivity and faster development time
  - Server infrastructure management is handled by someone else
- Easy to scale and horizontal scaling is managed by the platform
- Only pay for or consume resources when necessary and as needed
  - Never pay for idle resources
- Applications can be written in almost any programming language
- Built in availability and fault tolerance



## Disadvantages of Function as a Service

- Decreased transparency. Hard to understand the entire system.
  - Potentially tough to debug
- Auto-scaling of function calls often means auto-scaling of cost.
  - Makes it tough to estimate cost
- Can be tough to keep track of many functions.
- Stateful applications are harder to build
- Hard to pack dependencies!



# Disadvantages of Function as a Service

- **Time and Resource Limitations**
  - AWS
  - Azure
  - Google
- Cold start Lag
- Vendor Lock-In
- Different Programming Paradigm, requiring time to migrate



# FaaS vs Containers

- User's Applications
  - FaaS: User creates a program code , specifies needed dependencies and uploads it. The provider then provisions the computing environment
  - Container: User creates a container definition file that specifies apps and configurations. The container is initiated from the image created by the user.
- Uptime
  - FaaS: Only when the code is running
  - Container: When the container is running
- Payment for the resources
  - FaaS: Only when the code is running
  - Container: A server is needed to run the container



# FaaS vs Containers

- Runtime Limitations
  - FaaS: Limited
  - Containers: Unlimited
- Hosting specificity
  - FaaS: Better suited for microservices
  - Containers: Suited for any type of services
- Programming languages support
  - FaaS: Limited
  - Containers: Unlimited





# AWS Lambda

- AWS Lambda (2014) was first successful FaaS
- AWS Lambda supports many programming languages:
  - Python, Java, Node.js, Go, Ruby, and C# (.Net core)
  - Even Linux executables!
- Aims for millisecond startup latency
  - Caching containers will likely mean significant speed-up from recent use
  - Warm-boot



# AWS Lambda Pricing

- Pricing based on **number of requests** and their **duration**
  - Request cost is \$0.20 per million per month (US East)
  - ... but the first million requests per month are free
- Duration cost is \$0.0000166667 per **GB-second**
  - So involves both time and allocated memory you've chosen
  - ... but the first 400,000 GB-seconds per month are free
  - Memory allocation can be as low as 128 MB
  - (So the free tier will go a long way, for small-scale applications)
- Welcome to try it out!





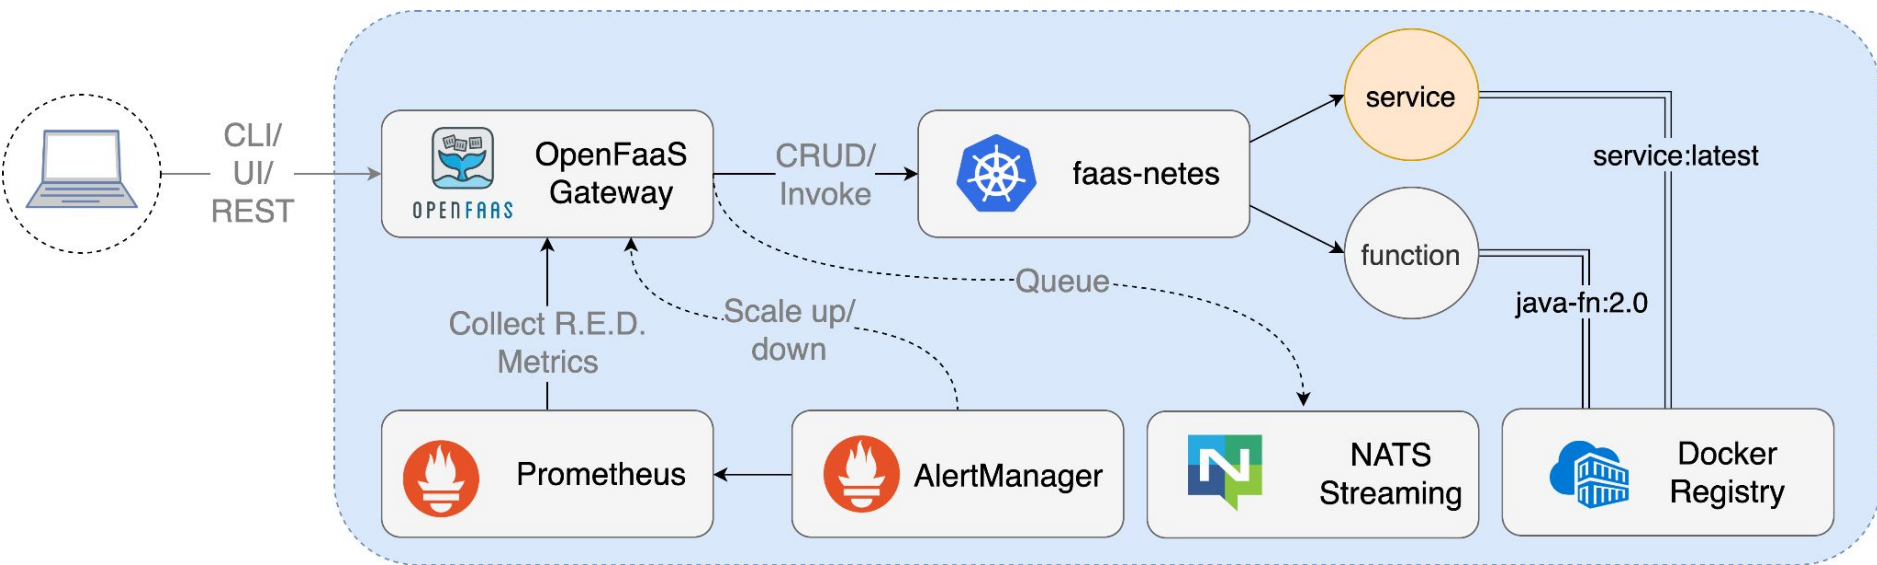
# OpenFaaS

# What is OpenFaaS?

- An open-source framework for building and running serverless applications on your own infrastructure
  - Public cloud providers do not disclose a lot about their infrastructure
- Originated from the serverless framework in Docker Swarm and now supports other kinds of infrastructure backends, such as **Kubernetes**
- **Functions in OpenFaaS are containers**
  - Any program written in any language can be packed as a function by leveraging the container technologies of Docker.

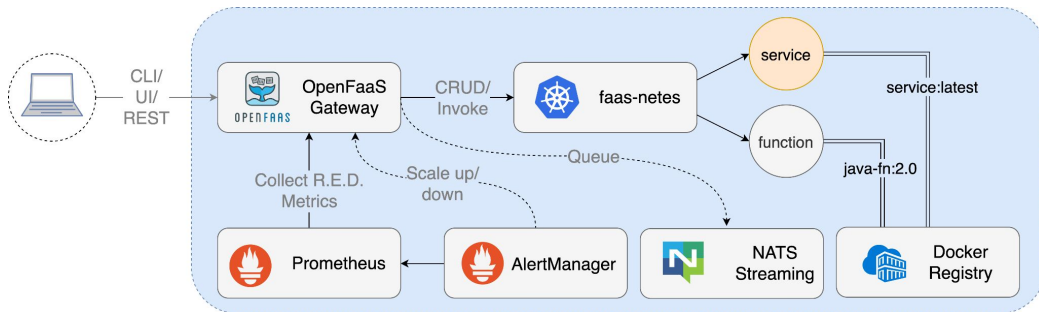


# OpenFaaS Infrastructure



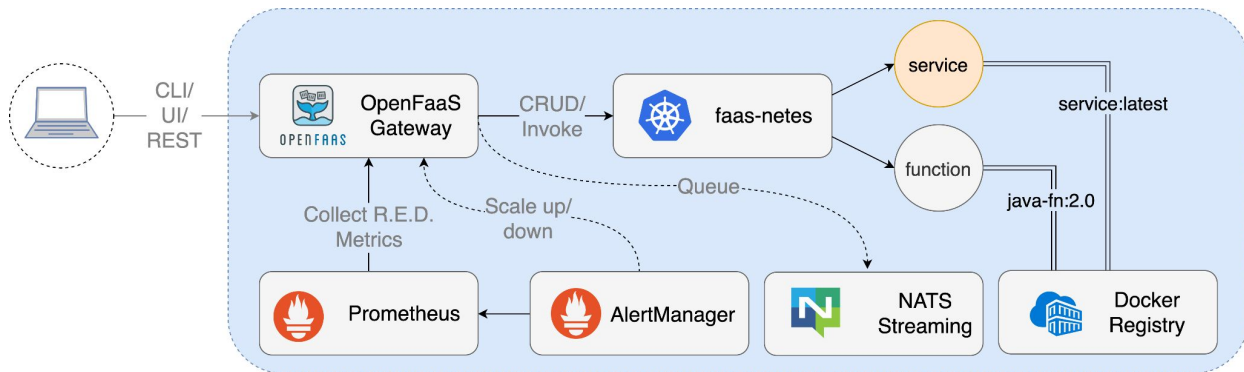
# OpenFaaS Infrastructure

- OpenFaaS consists of
  - An API gateway
  - CLI
  - Function watchdog
  - Prometheus
  - Kubernetes (or faasd) orchestration engines
  - Containers for Running the functions



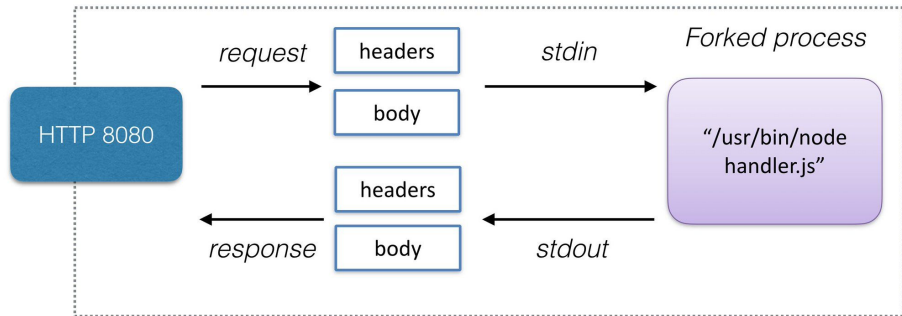
# OpenFaaS Infrastructure

- CLI: Either curl or faas-cli to connect to the API Gateway
- API Gateway: Receives user requests via the CLI or from web and invokes the functions via faas-netes (Kubernetes)



# OpenFaaS Infrastructure

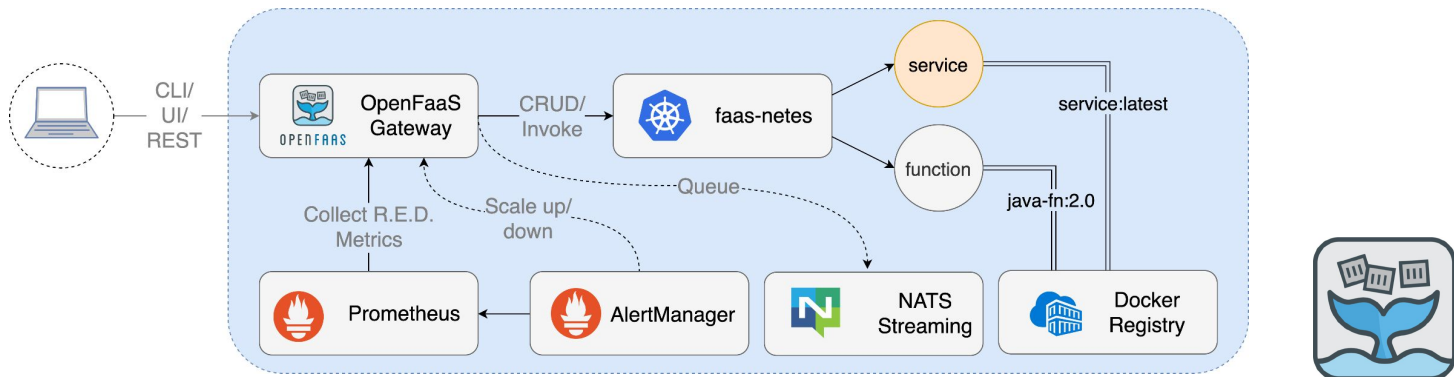
- Function Watchdog: Responsible for starting and monitoring functions
  - Any binary can become a function through the use of watchdog.
  - Requirement (now deprecated)
    - Accept input via the standard input (stdin) and print out the result to the standard output (stdout).
    - of-watchdog allows alternative to STDIO for communication between the watchdog and the function.





# OpenFaaS Infrastructure

- Prometheus: Separate monitoring and time-series database
  - <https://prometheus.io/>
- AlertManager: Reads usage (requests per second) metrics from Prometheus and fires alerts to the API Gateway for auto-scaling



# OpenFaaS Installation

- How to install OpenFaaS?
  - <https://docs.openfaas.com/cli/install/>
  - PLONK(F) Stack
    - Prometheus, **Linux**, OpenFaaS, NATs, faasd
  - <https://docs.openfaas.com/deployment/>
  - <https://github.com/openfaas/faasd/>
- DO NOT USE Native Mac or Windows
  - Use a Linux subsystem or a VM on QEMU or Virtualbox
- We will use multipass (also will be used for HW 2)
  - <https://multipass.run/>





# TODOs!

- Kubernetes Quiz!
- HW 2
- Study for Midterms!





# Agenda for Today

- Function as a Service
- OpenFaaS
- Readings
  - Recommended: CCSA Chapter 4
  - Optional:
    - <https://aws.amazon.com/solutions/case-studies/netflix-and-aws-lambda/>





# Questions?

